

CS145 Howework 1

****Important Note:**** HW1 is due on **11:59 PM PT, Oct 19 (Monday, Week 3)**. Please submit through GradeScope (you will receive an invite to Gradescope for CS145 Fall 2020.).

Print Out Your Name and UID

****Name:** Rithika Srinivasan, **UID:** 905125793******

Before You Start

You need to first create HW1 conda environment by the given `cs145hw1.yml` file, which provides the name and necessary packages for this tasks. If you have `conda` properly installed, you may create, activate or deactivate by the following commands:

```
conda env create -f cs145hw1.yml
conda activate hw1
conda deactivate
```

OR

```
conda env create --name NAMEOFOURCHOICE -f cs145hw1.yml
conda activate NAMEOFOURCHOICE
conda deactivate
```

To view the list of your environments, use the following command:

```
conda env list
```

More useful information about managing environments can be found [here](#).

You may also quickly review the usage of basic Python and Numpy package, if needed in coding for matrix operations.

In this notebook, you must not delete any code cells in this notebook. If you change any code outside the blocks that you are allowed to edit (between `START/END YOUR CODE HERE`), you need to highlight these changes. You may add some additional cells to help explain your results and observations.

```
In [5]: import numpy as np
import pandas as pd
import sys
import random as rd
import matplotlib.pyplot as plt
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:
`%reload_ext autoreload`

If you can successfully run the code above, there will be no problem for environment setting.

1. Linear regression

This workbook will walk you through a linear regression example.

```
In [6]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test')
# As a sanity check, we print out the size of the training data (1000, 100) and
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape: ', lm.train_y.shape)

Training data shape: (1000, 100)
Training labels shape: (1000,)
```

1.1 Closed form solution

In this section, complete the `getBeta` function in `linear_regression.py` which use the close for solution of $\hat{\beta}$.

Train you model by using `lm.train('0')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [7]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test')
training_error= 0
testing_error= 0
#####
# START YOUR CODE HERE #
#####
lm.normalize()
beta = lm.train('0')
predicted_train_y = lm.predict(lm.train_x, beta)
training_error = lm.compute_mse(predicted_train_y, lm.train_y)

predicted_test_y = lm.predict(lm.test_x, beta)
testing_error = lm.compute_mse(predicted_test_y, lm.test_y)

#####
# END YOUR CODE HERE #
#####
print('Training error is: ', training_error)
print('Testing error is: ', testing_error)

Learning Algorithm Type: 0
y shape
(1000,)
x shape
(1000, 101)
Beta: [ 5.23000000e-01 -3.95099505e-02 -3.01401932e-02 -5.71438644e-02
 -1.72769796e-02 -4.13700127e-03 -5.86318630e-02 -6.89027284e-02
```

```

-3.56331805e-02 -1.87845537e-02 -1.82888714e-02 5.29276130e-02
2.53519018e-02 -4.15812928e-02 -3.30193382e-02 2.65867992e-03
1.34068950e-02 -3.88013327e-02 4.11038867e-02 -2.32239983e-02
-2.68494719e-02 -5.67582270e-02 2.85948574e-02 -5.22058491e-02
-1.94232592e-02 4.61988692e-02 -3.87491283e-02 3.82055256e-02
1.27021593e-02 5.82271850e-02 -4.20937718e-02 -8.05582038e-02
5.50688227e-02 -2.88202457e-02 -1.94706479e-02 2.58596756e-03
2.55048685e-02 1.39991237e-02 -3.38312079e-02 -1.80218433e-02
-8.42135902e-03 -5.61252496e-02 -3.60939866e-02 -1.12787490e-04
-4.02969672e-02 -1.20851201e-02 -1.41809480e-02 5.11770552e-03
4.48842190e-02 1.42864924e-02 1.79066117e-02 -3.08841654e-02
-3.67139837e-02 3.83560781e-02 -4.47435146e-02 -6.08180754e-02
4.69774181e-02 -5.86346690e-02 1.62361334e-02 -6.06942237e-02
-3.38205570e-02 -4.24317897e-02 -5.46648364e-02 -2.89378305e-02
5.33687506e-02 -3.17462303e-02 2.12826319e-03 -3.26837546e-02
6.84819052e-03 1.25455103e-02 -4.09640271e-02 8.88512549e-03
1.94883628e-02 6.04797247e-02 -4.23185183e-02 -4.76582979e-02
-6.69833777e-02 5.66019062e-02 4.63178581e-03 4.13664903e-02
7.10828556e-02 4.08986579e-02 -6.46605942e-02 3.05062530e-02
6.11970818e-02 -6.13118531e-04 4.12093831e-02 8.04511196e-05
3.21203863e-02 5.30651849e-02 -2.83935172e-02 -4.22856651e-02
4.23271015e-02 -1.72635991e-03 -6.75124152e-02 -3.30151234e-02
-2.14687553e-03 -6.00152621e-02 4.30059659e-02 6.79904935e-02
-3.84367853e-03]

```

Training error is: 0.08693886675396784

Testing error is: 0.11017540281675804

1.2 Batch gradient descent

In this section, complete the `getBetaBatchGradient` function in `linear_regression.py` which compute the gradient of the objective function.

Train you model by using `lm.train('1')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```

In [8]: lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test')
training_error= 0
testing_error= 0
#=====#
# STRART YOUR CODE HERE #
#=====#
# lm.normalize()
beta2 = lm.train('1')
predicted_train_y = lm.predict(lm.train_x, beta2)
training_error = lm.compute_mse(predicted_train_y, lm.train_y)

predicted_test_y = lm.predict(lm.test_x, beta2)
testing_error = lm.compute_mse(predicted_test_y, lm.test_y)

#=====#
# END YOUR CODE HERE #
#=====#
print('Training accuracy is: ', training_error)
print('Testing accuracy is: ', testing_error)

```

Learning Algorithm Type: 1

Training accuracy is: 0.08694299599095222

Testing accuracy is: 0.11024614457982801

1.3 Stochastic gradient descent

In this section, complete the `getBetaStochasticGradient` function in `linear_regression.py`, which use an estimated gradient of the objective function.

Train you model by using `lm.train('2')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [9]: lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test')
training_error= 0
testing_error= 0
#####
# START YOUR CODE HERE #
#####
lm.normalize()
beta3 = lm.train('2')
print("beta shape: {}".format(beta3.shape))
predicted_train_y = lm.predict(lm.train_x, beta3)
training_error = lm.compute_mse(predicted_train_y, lm.train_y)

predicted_test_y = lm.predict(lm.test_x, beta3)
testing_error = lm.compute_mse(predicted_test_y, lm.test_y)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_error)
print('Testing accuracy is: ', testing_error)
```

Learning Algorithm Type: 2

beta: [5.27260629e-01 -3.22490083e-02 -2.52778364e-02 -7.06036388e-02

-2.82354510e-02 3.56668202e-02 -4.70891581e-02 -7.57249454e-02

-2.35042789e-02 2.81274688e-03 -4.46334966e-03 6.95732170e-02

2.43536507e-02 -5.23413242e-02 -1.89250970e-02 1.04886265e-02

-3.16493883e-03 -3.63267560e-02 4.29911807e-02 -2.95306646e-02

-1.49627856e-02 -7.07073352e-02 4.11758087e-02 -1.78392945e-02

-2.45146168e-02 7.94358034e-02 -3.84546251e-02 4.05421998e-02

3.51370715e-02 1.08159615e-01 -3.04196076e-02 -1.09225607e-01

4.18097838e-02 -2.14100077e-02 -2.09306075e-02 2.01561313e-02

2.31726355e-02 2.09414656e-02 -5.71102529e-02 3.43026875e-04

-1.50662680e-02 -7.01256778e-02 -2.35008893e-02 -1.53134156e-02

-3.74197261e-02 -1.51730935e-02 5.14151802e-04 -1.79295083e-02

5.07170746e-02 7.32814765e-03 2.52680257e-02 -3.44471047e-02

-4.41691822e-02 5.56941732e-02 -4.59119668e-02 -4.61664008e-02

5.50444992e-02 -6.34538920e-02 4.02294047e-02 -2.33090825e-02

-2.34403532e-02 -2.03412452e-02 -6.36716510e-02 -4.47915248e-02

6.92972447e-02 -6.44205665e-02 -2.58217647e-02 -3.33776821e-02

6.83578939e-03 2.79441471e-02 -6.38713402e-02 3.35247472e-03

1.16311083e-02 5.02485442e-02 -3.93100840e-02 -6.56327722e-02

-6.44397514e-02 4.40650187e-02 -1.03643931e-02 3.61635994e-02

9.81052000e-02 2.70771613e-02 -6.81477490e-02 5.67842272e-02

4.95537479e-02 7.87280014e-03 2.64716020e-02 3.97973853e-04

3.75332601e-02 5.89671590e-02 -3.37741464e-02 -5.44337262e-02

5.69887385e-02 -1.68717547e-02 -5.32183805e-02 -6.45138326e-02

-3.32494904e-02 -7.15423468e-02 3.39220730e-02 7.38992551e-02

1.20473436e-02]

Beta: [5.27260629e-01 -3.22490083e-02 -2.52778364e-02 -7.06036388e-02

-2.82354510e-02 3.56668202e-02 -4.70891581e-02 -7.57249454e-02

```

-2.35042789e-02  2.81274688e-03 -4.46334966e-03  6.95732170e-02
 2.43536507e-02 -5.23413242e-02 -1.89250970e-02  1.04886265e-02
-3.16493883e-03 -3.63267560e-02  4.29911807e-02 -2.95306646e-02
-1.49627856e-02 -7.07073352e-02  4.11758087e-02 -1.78392945e-02
-2.45146168e-02  7.94358034e-02 -3.84546251e-02  4.05421998e-02
 3.51370715e-02  1.08159615e-01 -3.04196076e-02 -1.09225607e-01
 4.18097838e-02 -2.14100077e-02 -2.09306075e-02  2.01561313e-02
 2.31726355e-02  2.09414656e-02 -5.71102529e-02  3.43026875e-04
-1.50662680e-02 -7.01256778e-02 -2.35008893e-02 -1.53134156e-02
-3.74197261e-02 -1.51730935e-02  5.14151802e-04 -1.79295083e-02
 5.07170746e-02  7.32814765e-03  2.52680257e-02 -3.44471047e-02
-4.41691822e-02  5.56941732e-02 -4.59119668e-02 -4.61664008e-02
 5.50444992e-02 -6.34538920e-02  4.02294047e-02 -2.33090825e-02
-2.34403532e-02 -2.03412452e-02 -6.36716510e-02 -4.47915248e-02
 6.92972447e-02 -6.44205665e-02 -2.58217647e-02 -3.33776821e-02
 6.83578939e-03  2.79441471e-02 -6.38713402e-02  3.35247472e-03
 1.16311083e-02  5.02485442e-02 -3.93100840e-02 -6.56327722e-02
-6.44397514e-02  4.40650187e-02 -1.03643931e-02  3.61635994e-02
 9.81052000e-02  2.70771613e-02 -6.81477490e-02  5.67842272e-02
 4.95537479e-02  7.87280014e-03  2.64716020e-02  3.97973853e-04
 3.75332601e-02  5.89671590e-02 -3.37741464e-02 -5.44337262e-02
 5.69887385e-02 -1.68717547e-02 -5.32183805e-02 -6.45138326e-02
-3.32494904e-02 -7.15423468e-02  3.39220730e-02  7.38992551e-02
 1.20473436e-02]

```

beta shape: (101,)

Training accuracy is: 0.11373741198992812

Testing accuracy is: 0.13756413236233234

Questions:

1. Compare the MSE on the testing dataset for each version. Are they the same? Why or why not?
2. Apply z-score normalization for each feature and comment whether or not it affects the three algorithms.
3. Ridge regression is adding an L2 regularization term to the original objective function of mean squared error. The objective function becomes following:

$$J(\beta) = \frac{1}{2n} \sum_i (x_i^T \beta - y_i)^2 + \frac{\lambda}{2n} \sum_j \beta_j^2,$$

where $\lambda \geq 0$, which is a hyper parameter that controls the trade off. Take the derivative of this provided objective function and derive the closed form solution for β .

[Please type your answer here!](#)

Your answer here:

1. The MSE for closed form solution is very similar to the descents, around 0.1, because they are essentially producing the same functions, but the batch solution's beta is more of an approximate of inverting the matrix, by calculating iteratively instead of all at once. The stochastic gradient descent example is based off of a random order of values with each time its function is called, so it changes every time it is called.
2. Adding z-score normalization does not affect the first, closed form feature, since the data values are not affected at each step. However, the z score normalization raises the MSE for the batch gradient, because it does not account for outliers that it does not map. Also, the

stochastic gradient example would not work without normalized data at all because the path the gradient follows would be impossible to discover.

$$\begin{aligned}
 J(\beta) &= \frac{1}{2n} \sum_i (x_i^T \beta - y_i)^2 + \underbrace{\frac{\lambda}{2n} \sum_j \beta_j^2}_{\text{L2 regularization term}} \\
 &\text{in matrix form:} \\
 J(\beta) &= \frac{(X\beta - y)^T (X\beta - y)}{2n} + \frac{\lambda}{2n} \beta^T \beta \\
 &= \frac{1}{2n} (\beta^T X^T - y^T) (X\beta - y) + \frac{\lambda}{2n} \beta^T \beta \\
 &= \frac{1}{2n} (\beta^T X^T X \beta - y^T X \beta - \beta^T X^T y + y^T y) + \frac{\lambda}{2n} \beta^T \beta \\
 \frac{dJ}{d\beta} &= \frac{1}{2n} (2 X^T X \beta - 2 X^T y) + \frac{2\lambda}{2n} \beta = 0 \\
 &= \frac{1}{n} (X^T X \beta - X^T y + \lambda \beta) = 0 \\
 &= \frac{1}{n} ((X^T X + \lambda) \beta - X^T y) = 0 \\
 &= \frac{X^T X + \lambda}{n} \beta - \frac{X^T y}{n} = 0 \\
 &= \frac{X^T X + \lambda}{n} \beta = \frac{X^T y}{n} \\
 \boxed{\beta} &= (X^T X + \lambda)^{-1} X^T y
 \end{aligned}$$

3.

2. Logistic regression

This workbook will walk you through a logistic regression example.

```
In [37]: from hw1code.logistic_regression import LogisticRegression

lm = LogisticRegression()
```

```
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-
# As a sanity check, we print out the size of the training data (1000, 5) and tr
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape: ', lm.train_y.shape)
```

```
Training data shape: (1000, 5)
Training labels shape: (1000,)
```

2.1 Batch gradiend descent

In this section, complete the `getBeta_BatchGradient` in `logistic_regression.py`, which compute the gradient of the log likelihood function.

Complete the `compute_avglogL` function in `logistic_regression.py` for sanity check.

Train you model by using `lm.train('0')` function.

And print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```
In [41]: lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####
lm.normalize()
beta4 = lm.train('0')

# predicted_train_y = lm.predict(lm.train_x, beta4)
# training_error = lm.compute_accuracy(predicted_train_y, lm.train_y)

# predicted_test_y = lm.predict(lm.test_x, beta4)
# testing_error = lm.compute_accuracy(predicted_test_y, lm.test_y)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)
```

```
average logL for iteration 0: -2.0960887433574817
```

```
/Users/rithika/Desktop/sad145/hw1/hw1code/logistic_regression.py:81: RuntimeWarn
ing: invalid value encountered in double_scalars
```

```
stupid = exponent_beta_T / (1 + exponent_beta_T)
```

```
average logL for iteration 1000: nan
```

```
average logL for iteration 2000: nan
```

```
average logL for iteration 3000: nan
```

```
average logL for iteration 4000: nan
```

```
average logL for iteration 5000: nan
```

```
average logL for iteration 6000: nan
```

```
average logL for iteration 7000: nan
```

```
average logL for iteration 8000: nan
```

```
average logL for iteration 9000: nan
```

```
Training avgLogL: nan
```

```
Training accuracy is: 0
```

```
Testing accuracy is: 0
```

2.2 Newton Raphhson

In this section, complete the `getBeta_Newton` in `logistic_regression.py`, which make use of both first and second derivative.

Train you model by using `lm.train('1')` function.

Print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```
In [ ]: lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####

#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)
```

Questions:

1. Compare the accuracy on the testing dataset for each version. Are they the same? Why or why not?
2. Regularization. Similar to linear regression, an regularization term could be added to logistic regression. The objective function becomes following:

$$J(\beta) = -\frac{1}{n} \sum_i (y_i x_i^T \beta - \log(1 + \exp\{x_i^T \beta\})) + \lambda \sum_j \beta_j^2,$$

where $\lambda \geq 0$, which is a hyper parameter that controls the trade off. Take the derivative $\frac{\partial J(\beta)}{\partial \beta_j}$ of this provided objective function and provide the batch gradient descent update.

Your answer here:

[Please type your answer here!](#)

2.3 Visualize the decision boundary on a toy dataset

In this subsection, you will use the same implementation for another small dataset with each datapoint x with only two features (x_1, x_2) to visualize the decision boundary of logistic regression model.

```
In [ ]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression(verbose = False)
lm.load_data('./data/logistic-regression-toy.csv', './data/logistic-regression-to
# As a sanity chech, we print out the size of the training data (99,2) and train
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)
```


In the following block, you can apply the same implementation of logistic regression model (either in 2.1 or 2.2) to the toy dataset. Print out the $\hat{\beta}$ after training and accuracy on the train set.

```
In [ ]: training_accuracy= 0
#=====#
#  STRART YOUR CODE HERE  #
#=====#

#=====#
#   END YOUR CODE HERE   #
#=====#
print('Training accuracy is: ', training_accuracy)
```

Next, we try to plot the decision boundary of your learned logistic regression classifier.

Generally, a decision boundary is the region of a space in which the output label of a classifier is ambiguous. That is, in the given toy data, given a datapoint $x = (x_1, x_2)$ on the decision boundary, the logistic regression classifier cannot decide whether $y = 0$ or $y = 1$.

Question

Is the decision boundary for logistic regression linear? Why or why not?

Your answer here:

[Please type your answer here!](#)

Draw the decision boundary in the following cell. Note that the code to plot the raw data points are given. You may need `plt.plot` function (see [here](#)).

```
In [ ]: # scatter plot the raw data
df = pd.concat([lm.train_x, lm.train_y], axis=1)
groups = df.groupby("y")
for name, group in groups:
    plt.plot(group["x1"], group["x2"], marker="o", linestyle="", label=name)

# plot the decision boundary on top of the scattered points
#=====#
#  STRART YOUR CODE HERE  #
#=====#

#=====#
#   END YOUR CODE HERE   #
#=====#
plt.show()
```

End of Homework 1 :)

After you've finished the homework, please print out the entire `ipynb` notebook and two `py` files into one PDF file. Make sure you include the output of code cells and answers for questions. Prepare submit it to GradeScope.

```

import pandas as pd
import numpy as np
import sys
import random as rd

#insert an all-one column as the first column
def addAllOneColumn(matrix):
    n = matrix.shape[0] #total of data points
    p = matrix.shape[1] #total number of attributes

    newMatrix = np.zeros((n,p+1))
    newMatrix[:,1:] = matrix
    newMatrix[:,0] = np.ones(n)

    return newMatrix

# Reads the data from CSV files, converts it into Dataframe and returns x
and y dataframes
def getDataframe(filePath):
    dataframe = pd.read_csv(filePath)
    y = dataframe['y']
    x = dataframe.drop('y', axis=1)
    return x, y

# train_x and train_y are numpy arrays
# function returns value of beta calculated using (0) the formula beta =

$$(X^T X)^{-1} * (X^T Y)$$

def getBeta(train_x, train_y):
    n = train_x.shape[0] #total of data points
    p = train_x.shape[1] #total number of attributes

    beta = np.zeros(p)

    #=====#
    # STRART YOUR CODE HERE #
    #=====#
    Xt = np.transpose(train_x)
    X = train_x
    Y = train_y
    print("y shape")
    print(train_y.shape)

    print("x shape")
    print(train_x.shape)

    beta = np.matmul((np.linalg.inv(np.matmul(Xt,X))), (np.matmul(Xt,Y)))

#     print(beta)

    #=====#
    #     END YOUR CODE HERE #
    #=====#

```

```

    return beta

# train_x and train_y are numpy arrays
# lr (learning rate) is a scalar
# function returns value of beta calculated using (1) batch gradient
descent
def getBetaBatchGradient(train_x, train_y, lr, num_iter):
    beta = np.random.rand(train_x.shape[1])

    n = train_x.shape[0] #total of data points
    p = train_x.shape[1] #total number of attributes

    beta = np.random.rand(p)
    #update beta iteratively
    for iter in range(0, num_iter):
        deriv = np.zeros(p)
        for i in range(n):
            #=====#
            # STRART YOUR CODE HERE  #
            #=====#
            xi = train_x[i]
            xiT = np.transpose(xi)
            yi = train_y[i]

            #first_beta = (np.matmul(xiT,beta)
            first_beta = xiT.dot(beta)
            subtract_dat = np.subtract(first_beta, yi)

            deriv += xi.dot(subtract_dat)
            #=====#
            #   END YOUR CODE HERE   #
            #=====#
        deriv = deriv / n
        beta = beta - deriv.dot(lr)
    return beta

# train_x and train_y are numpy arrays
# lr (learning rate) is a scalar
# function returns value of beta calculated using (2) stochastic gradient
descent
def getBetaStochasticGradient(train_x, train_y, lr):
    n = train_x.shape[0] #total of data points
    p = train_x.shape[1] #total number of attributes

    beta = np.random.rand(p)

    epoch = 100
    for iter in range(epoch):
        indices = list(range(n))
        rd.shuffle(indices)
        for i in range(n):
            idx = indices[i]
            #=====#

```

```

        # STRART YOUR CODE HERE #
        #=====#
        xi = train_x[idx]
        xiT = np.transpose(xi)
        first_op = np.matmul(xiT, beta)

        yi = train_y[idx]
        second_op = yi - first_op
        third_op = lr * second_op * xi
        beta += third_op

        #=====#
        #   END YOUR CODE HERE   #
        #=====#

#         beta +=
    print("beta: {}".format(beta))
    return beta

# Linear Regression implementation
class LinearRegression(object):
    # Initializes by reading data, setting hyper-parameters, and forming
    linear model
    # Forms a linear model (learns the parameter) according to type of
    beta (0 - closed form, 1 - batch gradient, 2 - stochastic gradient)
    # Performs z-score normalization if z_score is 1
    def __init__(self, lr=0.005, num_iter=1000):
        self.lr = lr
        self.num_iter = num_iter
        self.train_x = pd.DataFrame()
        self.train_y = pd.DataFrame()
        self.test_x = pd.DataFrame()
        self.test_y = pd.DataFrame()
        self.algType = 0
        self.isNormalized = 0

    def load_data(self, train_file, test_file):
        self.train_x, self.train_y = getDataframe(train_file)
        self.test_x, self.test_y = getDataframe(test_file)

    def normalize(self):
        # Applies z-score normalization to the dataframe and returns a
        normalized dataframe
        self.isNormalized = 1
        means = self.train_x.mean(0)
        std = self.train_x.std(0)
        self.train_x = (self.train_x - means).div(std)
        self.test_x = (self.test_x - means).div(std)

    # Gets the beta according to input
    def train(self, algType):
        self.algType = algType

```

```

        newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-
one column as the first column
        print('Learning Algorithm Type: ', algType)
        if(algType == '0'):
            beta = getBeta(newTrain_x, self.train_y.values)
            print('Beta: ', beta)

            elif(algType == '1'):
                beta = getBetaBatchGradient(newTrain_x, self.train_y.values,
self.lr, self.num_iter)
                #print('Beta: ', beta)
            elif(algType == '2'):
                beta = getBetaStochasticGradient(newTrain_x,
self.train_y.values, self.lr)
                print('Beta: ', beta)
            else:
                print('Incorrect beta_type! Usage: 0 - closed form solution,
1 - batch gradient descent, 2 - stochastic gradient descent')

        return beta

# Predicts the y values on given data and learned beta
def predict(self,x, beta):
    newTest_x = addAllOneColumn(x)
    self.predicted_y = newTest_x.dot(beta)
    return self.predicted_y

# predicted_y and y are the predicted and actual y values
respectively as numpy arrays
# function returns the mean squared error (MSE) value for the test
dataset
def compute_mse(self,predicted_y, y):
    mse = np.sum((predicted_y - y)**2)/predicted_y.shape[0]
    return mse

```

```

# -*- coding: utf-8 -*-

import pandas as pd
import numpy as np
import sys
import random as rd

#insert an all-one column as the first column
def addAllOneColumn(matrix):
    n = matrix.shape[0] #total of data points
    p = matrix.shape[1] #total number of attributes

    newMatrix = np.zeros((n,p+1))
    newMatrix[:,0] = np.ones(n)
    newMatrix[:,1:] = matrix

    return newMatrix

# Reads the data from CSV files, converts it into Dataframe and returns x
and y dataframes
def getDataframe(filePath):
    dataframe = pd.read_csv(filePath)
    y = dataframe['y']
    x = dataframe.drop('y', axis=1)
    return x, y

# sigmoid function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# compute average logL
def compute_avglogL(X,y,beta):
    eps = 1e-50
    n = y.shape[0]
    avglogL = 0
    #=====#
    # STRART YOUR CODE HERE #
    #=====#

    for i in range(n):
        xT = X[i].T
        xT_beta = xT.dot(beta)
        y_xT_beta = y[i] * xT_beta

        exponent = np.exp(xT_beta)
        logarithika = np.log(1 + exponent)
        avglogL += y_xT_beta - logarithika
    avglogL = avglogL/n
    #=====#
    # END YOUR CODE HERE #
    #=====#
    return avglogL

```

```

# train_x and train_y are numpy arrays
# lr (learning rate) is a scalar
# function returns value of beta calculated using (0) batch gradient
descent
def getBeta_BatchGradient(train_x, train_y, lr, num_iter, verbose):
    beta = np.random.rand(train_x.shape[1])

    n = train_x.shape[0] #total of data points
    p = train_x.shape[1] #total number of attributes

    beta = np.random.rand(p)

    #update beta iteratively
    for iter in range(0, num_iter):
        #=====#
        # STRART YOUR CODE HERE #
        #=====#
        deriv_logL_beta = np.zeros(p)
        for i in range(n):
            yi = train_y[i] #()
            beta_T = np.transpose(beta) #(6,)
            xi = train_x[i] #(6,)
            thing_in_exp = np.matmul(beta_T, xi) #()
            exponent_beta_T = np.exp(thing_in_exp) #()

            stupid = exponent_beta_T / (1 + exponent_beta_T)
            inner = yi - stupid

            smthing_else_for_now = xi * stupid
            print("shape of stupid: {}".format(exponent_beta_T.shape))
            deriv_logL_beta = deriv_logL_beta + smthing_else_for_now
            beta = beta + deriv_logL_beta.dot(lr)

        #
        #
        #
        #
        #
        print("shape of yi: {}".format(yi.shape))
        #
        print("shape of xij: {}".format(xij.shape))
        #
        print("shape of beta_T: {}".format(beta_T.shape))
        #
        print("shape of xi: {}".format(xi.shape))
        #
        print("shape of exponent_beta_T:
        {}".format(exponent_beta_T.shape))

        #=====#
        # END YOUR CODE HERE #
        #=====#
        if(verbose == True and iter % 1000 == 0):
            avgLogL = compute_avglogL(train_x, train_y, beta)
            print(f'average logL for iteration {iter}: {avgLogL} \t')
    return beta

# train_x and train_y are numpy arrays
# function returns value of beta calculated using (1) Newton-Raphson
method

```



```

def getBeta_Newton(train_x, train_y, num_iter, verbose):
    n = train_x.shape[0] #total of data points
    p = train_x.shape[1] #total number of attributes

    beta = np.random.rand(p)
    for iter in range(0, num_iter):
        =====#
        # STRART YOUR CODE HERE #
        =====#

        =====#
        # END YOUR CODE HERE #
        =====#
        if(verbose == True and iter % 500 == 0):
            avgLogL = compute_avglogL(train_x, train_y, beta)
            print(f'average logL for iteration {iter}: {avgLogL} \t')
    return beta

# Logistic Regression implementation
class LogisticRegression(object):
    # Initializes by reading data, setting hyper-parameters
    # Learns the parameter using (0) Batch gradient (1) Newton-Raphson
    # Performs z-score normalization if isNormalized is 1
    # Print intermidate training loss if verbose = True
    def __init__(self,lr=0.005, num_iter=10000, verbose = True):
        self.lr = lr
        self.num_iter = num_iter
        self.verbose = verbose
        self.train_x = pd.DataFrame()
        self.train_y = pd.DataFrame()
        self.test_x = pd.DataFrame()
        self.test_y = pd.DataFrame()
        self.algType = 0
        self.isNormalized = 0

    def load_data(self, train_file, test_file):
        self.train_x, self.train_y = getDataframe(train_file)
        self.test_x, self.test_y = getDataframe(test_file)

    def normalize(self):
        # Applies z-score normalization to the dataframe and returns a
normalized dataframe
        self.isNormalized = 1
        data = np.append(self.train_x, self.test_x, axis = 0)
        means = data.mean(0)
        std = data.std(0)
        self.train_x = (self.train_x - means).div(std)
        self.test_x = (self.test_x - means).div(std)

    # Gets the beta according to input
    def train(self, algType):

```

```

        self.algType = algType
        newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-
one column as the first column
        if(algType == '0'):
            beta = getBeta_BatchGradient(newTrain_x, self.train_y.values,
self.lr, self.num_iter, self.verbose)
            #print('Beta: ', beta)

            elif(algType == '1'):
                beta = getBeta_Newton(newTrain_x, self.train_y.values,
self.num_iter, self.verbose)
                #print('Beta: ', beta)
            else:
                print('Incorrect beta_type! Usage: 0 - batch gradient
descent, 1 - Newton-Raphson method')

        train_avglogL = compute_avglogL(newTrain_x, self.train_y.values,
beta)
        print('Training avgLogL: ', train_avglogL)

        return beta

# Predict on given data x with learned parameter beta
def predict(self, x, beta):
    newTest_x = addAllOneColumn(x)
    self.predicted_y = (sigmoid(newTest_x.dot(beta))>=0.5)
    return self.predicted_y

# predicted_y and y are the predicted and actual y values
respectively as numpy arrays
# function returns the accuracy
def compute_accuracy(self,predicted_y, y):
    acc = np.sum(predicted_y == y)/predicted_y.shape[0]
    return acc

```